

# White\_MLR\_Initial.R

nebojsahrnjez

2021-12-03

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --

## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.6      v stringr 1.4.0
## v tidyr   1.1.4      v forcats 0.5.1
## v readr   2.1.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(gridExtra)
```

```
##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##   combine
```

```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
library(leaps)
library(glmnet)
```

```
## Loading required package: Matrix
```

```
##
```

```
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':
```

```
##
```

```
##      expand, pack, unpack
```

```
## Loaded glmnet 4.1-3
```

```
library(coefplot)
library(ggfortify)
```

```
## Registered S3 methods overwritten by 'ggfortify':
```

```
##   method      from
```

```
##   autoplot.acf  useful
```

```
##   fortify.acf   useful
```

```
##   fortify.kmeans useful
```

```
##   fortify.ts    useful
```

```
library(readr)
library(car)
```

```
## Loading required package: carData
```

```
##
```

```
## Attaching package: 'car'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
##      some
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      recode
```

```
library(moments)
library(ggpubr)
library(ggrepel)
library(qqplotr)
```

```
##
```

```
## Attaching package: 'qqplotr'
```

```
## The following objects are masked from 'package:ggplot2':
```

```
##
```

```
##      stat_qq_line, StatQqLine
```

```
library(MASS)
```

```
##  
## Attaching package: 'MASS'  
  
## The following object is masked from 'package:dplyr':  
##  
##     select
```

```
library(ordinal)
```

```
##  
## Attaching package: 'ordinal'  
  
## The following object is masked from 'package:dplyr':  
##  
##     slice
```

```
library(caret)
```

```
## Loading required package: lattice  
  
##  
## Attaching package: 'caret'  
  
## The following object is masked from 'package:purrr':  
##  
##     lift
```

```
library(rpart)  
library(rpart.plot)  
library(rpartScore)  
library(DMwR2)
```

```
## Registered S3 method overwritten by 'quantmod':  
##   method             from  
##   as.zoo.data.frame zoo
```

```
library(randomForest)
```

```
## randomForest 4.6-14  
  
## Type rfNews() to see new features/changes/bug fixes.  
  
##  
## Attaching package: 'randomForest'
```

```

## The following object is masked from 'package:gridExtra':
##
##      combine

## The following object is masked from 'package:ggplot2':
##
##      margin

## The following object is masked from 'package:dplyr':
##
##      combine

white <- read_csv("winequality-white.csv")

## Rows: 4898 Columns: 12

## -- Column specification -----
## Delimiter: ","
## dbl (12): fixed acidity, volatile acidity, citric acid, residual sugar, chlo...

##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

sum(is.na(white))

## [1] 0

white <- na.omit(white)

dfw <- as.data.frame(white)

SEED <- 5864
set.seed(SEED)

test <- sample(1:nrow(dfw), size = nrow(dfw)/5)
train <- (-test)

dfw.train <- dfw[train,]
dfw.test <- dfw[test,]

#Create a test and training data set, 20/80 split

w.mlm.train <- lm(
  quality ~.,
  data = dfw.train
)
#Create linear regression with all predictors using the training dataset

```

```
w.mlm.predict <- predict(w.mlm.train, newdata = dfw.test)
w.mlm.predict.rounded <- round(w.mlm.predict, digits = 0)

#Using the training model, predict the response variable using the "test set"

w.mlm.predict.rounded <- round(w.mlm.predict, digits = 0)

#Round the predicted to a integer so it can be compared to the test set for
# classification

(con_mat <- table(w.mlm.predict.rounded, dfw.test$quality))
```

```
##
## w.mlm.predict.rounded   3   4   5   6   7   8   9
##                        4   0   3   0   0   0   0
##                        5   2  17 117  52   2   0
##                        6   4  12 175 336 141  21
##                        7   0   0   1  44  38  13   1
```

```
mean(w.mlm.predict.rounded==dfw.test$quality)
```

```
## [1] 0.5045965
```

```
mean(w.mlm.predict.rounded!=dfw.test$quality)
```


```
## [1] 0.4954035
```

```
#Create the confusion matrix and calculate the proportion correct and incorrect
#This model was correct ~52.91113% of the time
```

```
correctness <- rep(NA, 100)

for(j in 1:100){
  tol_value <- j*0.01
  tol_x <- w.mlm.predict+(w.mlm.predict*tol_value)
  tol_z <- w.mlm.predict-(w.mlm.predict*tol_value)
  yorn <- 0
  for(i in 1:nrow(dfw.test)){
    yorn[i]<-between(dfw.test$quality[i],tol_z[i],tol_x[i])
  }
  yorn
  correctness[j] <- sum(yorn)/length(yorn)
}
correctness <- correctness*100
tolerance_percentage <- 1:100
plot(tolerance_percentage,correctness,
     main = "Classification Rate vs Tolerance")
```

## Classification Rate vs Tolerance



*#This is basically MSE proof, at .50 so withing a range of 1 we have a 95%  
#confidence interval assuming confidence interval*

```
w.glm.train <- glm(quality~.,data=dfw.train, family = "poisson")
w.glm.predict <- predict(w.glm.train, newdata = dfw.test, type = "response")

w.glm.predict.rounded <- round(w.glm.predict, digits = 0)

con_mat.glm <- table(w.glm.predict.rounded, dfw.test$quality)
mean(w.glm.predict.rounded==dfw.test$quality)
```

```
## [1] 0.5097038
```

```
mean(w.glm.predict.rounded!=dfw.test$quality)
```

```
## [1] 0.4902962
```

*#Tried poisson regression, it is only marginally better at 50.97%*

*#Use a validation set to choose among models*

```
w.best <- regsubsets(quality~., data = dfw.train, nvmax=11)
#Create a best subsets model with all the training data and variables

test.mat=model.matrix(quality~., data = dfw.test)
```

```
w.val.errors <- rep(NA,11)
classif <- rep(NA,11)
inclass <- rep(NA,11)
for(i in 1:11){
  coefi <- coef(w.best, id=i)
  pred <- test.mat[,names(coefi)]%*%coefi
  w.val.errors[i] <- mean((dfw.test$quality-round(pred,digits=0))^2)
  classif[i] <- mean(round(pred,digits=0)==dfw.test$quality)
  inclass[i] <- mean(round(pred,digits=0)!=dfw.test$quality)
}
w.val.errors
```

```
## [1] 0.7568948 0.6762002 0.6537283 0.6843718 0.6608784 0.6547497 0.6618999
## [8] 0.6455567 0.6475996 0.6496425 0.6475996
```

```
classif
```

```
## [1] 0.4923391 0.5229826 0.5097038 0.4984678 0.5005107 0.4974464 0.4994893
## [8] 0.5066394 0.5076609 0.5056180 0.5045965
```

```
which.min(w.val.errors)
```

```
## [1] 8
```

```
which.max(classif)
```

```
## [1] 2
```

```
round(coef(w.best, 9),3)
```

```
##          (Intercept)          'fixed acidity'          'volatile acidity'
##          139.311          0.053          -1.921
##          'residual sugar' 'free sulfur dioxide' 'total sulfur dioxide'
##          0.078          0.004          0.000
##          density          pH          sulphates
##          -138.892          0.553          0.592
##          alcohol
##          0.205
```

```
round(coef(w.best, 6),3)
```

```
##          (Intercept)          'volatile acidity'          'residual sugar'
##          94.456          -1.989          0.058
##          'free sulfur dioxide'          density          sulphates
##          0.004          -92.157          0.576
##          alcohol
##          0.266
```

```
w.best2 <- regsubsets(quality~., data = dfw, nvmax=11)
round(coef(w.best2, 9),3)
```

```
##          (Intercept)          'fixed acidity'          'volatile acidity'
##          151.256          0.068          -1.872
##          'residual sugar' 'free sulfur dioxide' 'total sulfur dioxide'
##          0.082          0.004          0.000
##          density          pH          sulphates
##          -151.398          0.692          0.634
##          alcohol
##          0.194
```

```
round(coef(w.best2, 6),3)
```

```
##          (Intercept) 'volatile acidity' 'residual sugar'          density
##          116.300          -1.992          0.071          -115.304
##          pH          sulphates          alcohol
##          0.490          0.605          0.232
```

*#Found the model that has the highest classification rate, val.error doesnt tell  
# us much here other than if the prediction errors are normally distributed  
# then 95% of the time it will be within 1 of the correct value*

*#Use cross validation to choose among models*

```
predict.regsubsets <- function(object, newdata, id, ...){
  form <- as.formula(object$call[[2]])
  mat <- model.matrix(form,newdata)
  coefi <- coef(object, id=id)
  xvars <- names(coefi)
  mat[,xvars] %*% coefi
}

k <- 10
n <- nrow(dfw)
set.seed(SEED)
folds=sample(rep(1:k, length=n))

w.cv.errors=matrix(NA,k,11, dimnames=list(NULL, paste(1:11)))
classif.kfold <- matrix(NA,k,11, dimnames=list(NULL, paste(1:11)))

for(j in 1:k){
  best.fit=regsubsets(quality~.,data=dfw[folds!=j,],nvmax=11)
  for(i in 1:11){
    pred=predict(best.fit,dfw[folds==j,], id=i)
    w.cv.errors[j,i] <- mean( (dfw$quality[folds==j] - round(pred,digits =0))^2)
    classif.kfold[j,i] <- mean(round(pred,digits=0)==dfw$quality[folds==j])
  }
}

mean.w.cv.errors <- rep(0, 11)
for (i in 1:11) {
```



```

mean.w.cv.errors[i] <- mean(w.cv.errors[,i])
}

mean.classif.kfold <- rep(0, 11)
for (i in 1:11) {
  mean.classif.kfold[i] <- mean(classif.kfold[,i])
}

#Best-plot function
best.plot <- function(varName, varLabel, minmax=" ") {
  gg <- ggplot(data.frame(varName), aes(x=seq_along(varName),
                                         y=varName)) +

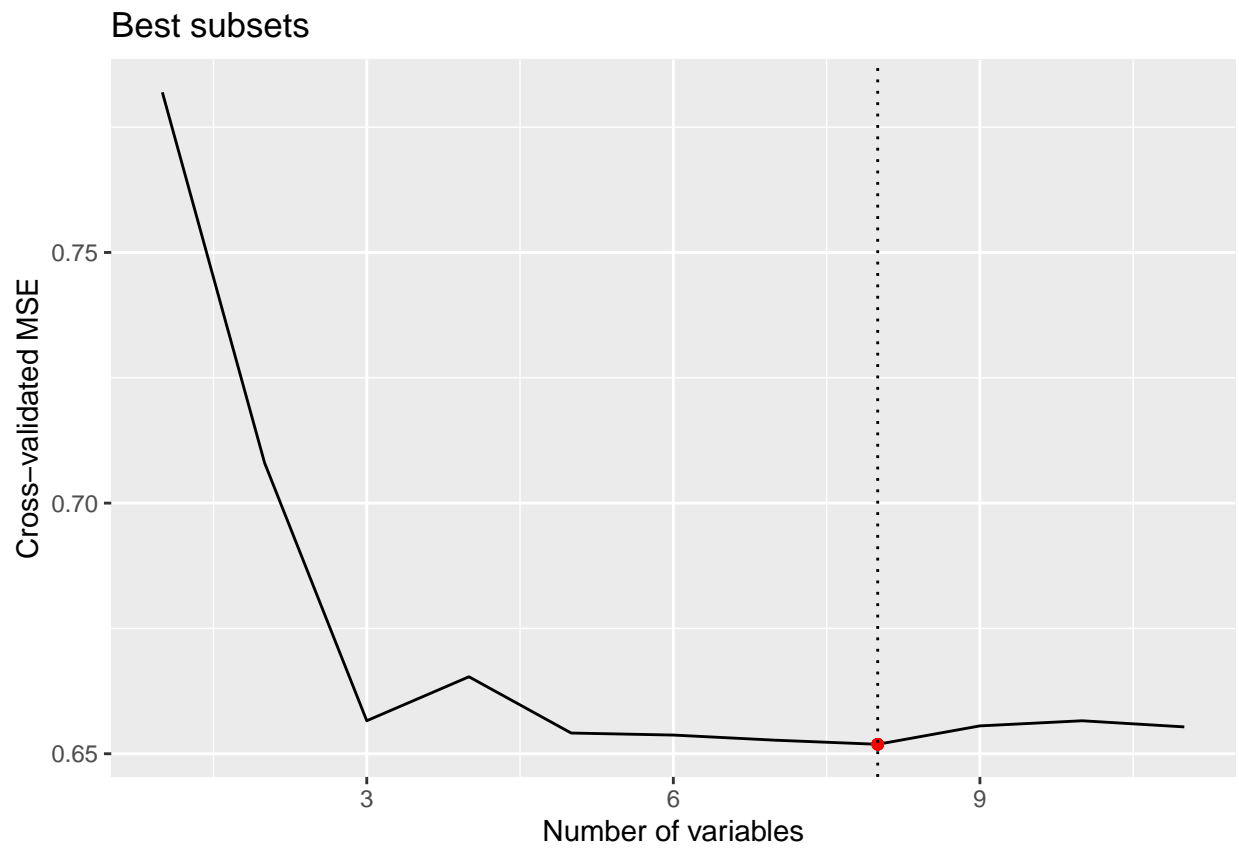
    geom_line() +
    labs(x="Number of variables",
         y=varLabel, title="Best subsets")

  if (minmax=="min") {
    gg <- gg + geom_point(aes(x=which.min(varName), y=min(varName)),
                          color="red") +
      geom_vline(aes(xintercept=which.min(varName)), linetype="dotted")
  }
  if (minmax=="max") {
    gg <- gg + geom_point(aes(x=which.max(varName), y=max(varName)),
                          color="red") +
      geom_vline(aes(xintercept=which.max(varName)), linetype="dotted")
  }
  return(gg)
}

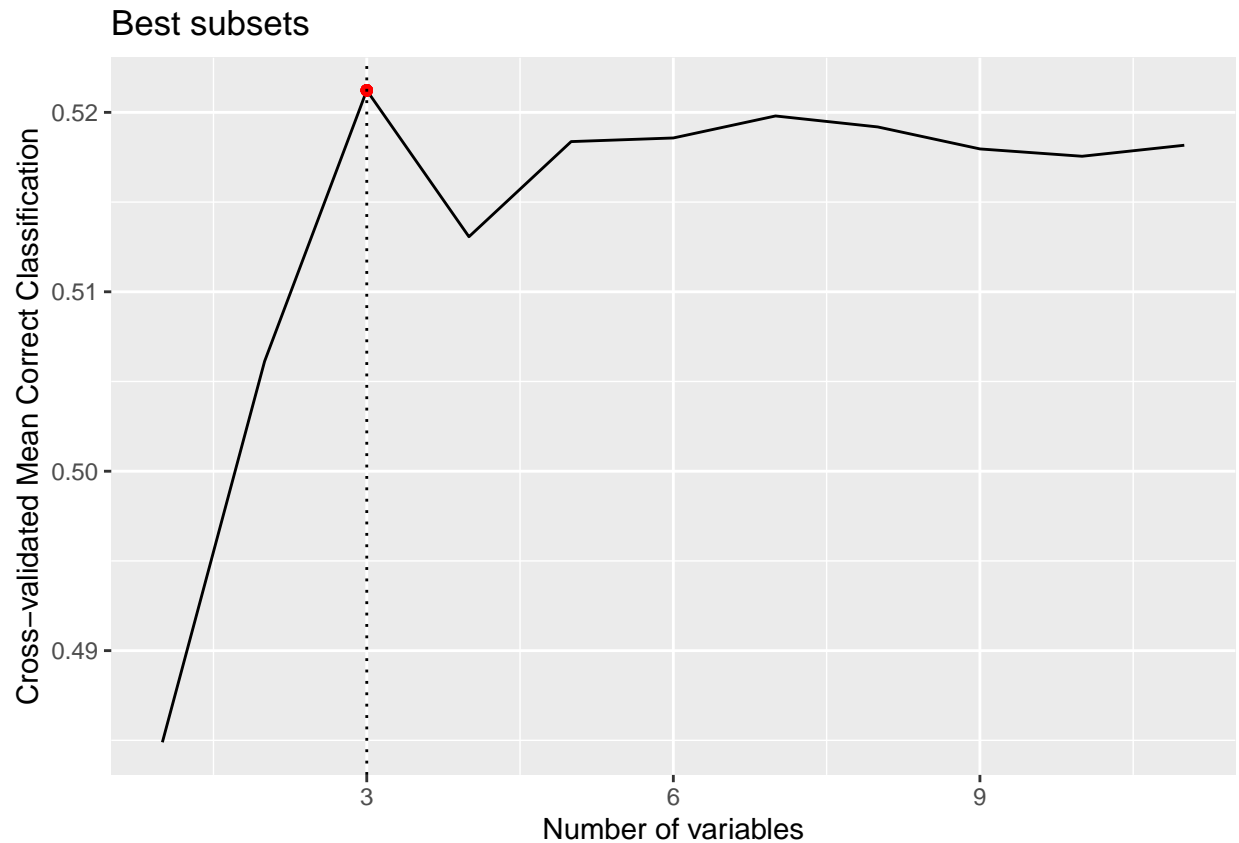
#End of best-plot function

best.plot(mean.w.cv.errors, "Cross-validated MSE", "min")

```



```
best.plot(mean.classif.kfold, "Cross-validated Mean Correct Classification",  
          "max")
```



```
mean.w.cv.errors[3]
```

```
## [1] 0.6565882
```

```
mean.classif.kfold
```

```
## [1] 0.4848913 0.5061229 0.5212278 0.5130675 0.5183707 0.5185735 0.5197997
## [8] 0.5191878 0.5179625 0.5175539 0.5181666
```

```
which.max(mean.classif.kfold)
```

```
## [1] 3
```

```
w.best2 <- regsubsets(quality~., data = dfw, nvmax=11)
round(coef(w.best2, 3),3)
```

```
##          (Intercept) 'volatile acidity'  'residual sugar'      alcohol
##              2.356             -2.107              0.027           0.375
```

```
#Use lasso regression as a potential for variable selection
```

```
library(dplyr)
library(tidyverse)
```

```

library(gridExtra)
library(corrplot)
library(leaps)
library(glmnet)
library(coefplot)
library(ggfortify)
library(readr)
library(car)
library(moments)
library(ggpubr)
library(ggrepel)
library(qqplotr)
library(MASS)
library(ordinal)
library(caret)
library(rpart)
library(rpart.plot)
library(rpartScore)
library(DMwR2)
library(randomForest)

white <- read_csv("winequality-white.csv")

```

```
## Rows: 4898 Columns: 12
```

```

## -- Column specification -----
## Delimiter: ","
## dbl (12): fixed acidity, volatile acidity, citric acid, residual sugar, chlo...

##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

```

```
sum(is.na(white))
```

```
## [1] 0
```

```

white <- na.omit(white)

dfw <- as.data.frame(white)

dfw$quality <- as.factor(dfw$quality)

SEED <- 5864
set.seed(SEED)

test <- sample(1:nrow(dfw), size = nrow(dfw)/5)
train <- (-test)

dfw.train <- dfw[train,]

```

```
dfw.test <- dfw[test,]
```

```
w.clm <- clm(quality~.,data=dfw.train)
```

```
## Warning: (3) Model is nearly unidentifiable: large eigenvalue ratio
## - Rescale variables?
## In addition: Absolute and relative convergence criteria were met
```

```
w.clm.predict <- predict(w.clm, newdata=dfw.test, type = "class")
```

```
w.clm.predict <- as.numeric(as.character(unlist(w.clm.predict)))
```

```
dfw.test$quality <- as.numeric(as.character(unlist(dfw.test$quality)))
```

```
mean(w.clm.predict==dfw.test$quality)
```

```
## [1] 0.5127681
```

```
summary(w.clm)
```

```
## formula:
```

```
## quality ~ 'fixed acidity' + 'volatile acidity' + 'citric acid' + 'residual sugar' + chlorides + 'free
```

```
## data: dfw.train
```

```
##
```

```
## link threshold nobs logLik AIC niter max.grad cond.H
```

```
## logit flexible 3919 -4366.67 8767.33 8(0) 9.08e-10 7.6e+11
```

```
##
```

```
## Coefficients:
```

```
## Estimate Std. Error z value Pr(>|z|)
```

```
## 'fixed acidity' 1.838e-01 6.626e-02 2.773 0.00555 **
```

```
## 'volatile acidity' -5.132e+00 3.445e-01 -14.897 < 2e-16 ***
```

```
## 'citric acid' 1.618e-01 2.728e-01 0.593 0.55314
```

```
## 'residual sugar' 2.180e-01 2.547e-02 8.561 < 2e-16 ***
```

```
## chlorides -1.039e+00 1.528e+00 -0.680 0.49640
```

```
## 'free sulfur dioxide' 1.300e-02 2.538e-03 5.122 3.02e-07 ***
```

```
## 'total sulfur dioxide' -7.160e-04 1.106e-03 -0.647 0.51748
```

```
## density -4.270e+02 6.817e+01 -6.264 3.76e-10 ***
```

```
## pH 1.642e+00 3.268e-01 5.025 5.04e-07 ***
```

```
## sulphates 1.724e+00 2.917e-01 5.910 3.42e-09 ***
```

```
## alcohol 4.620e-01 8.473e-02 5.453 4.95e-08 ***
```

```
## ---
```

```
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Threshold coefficients:
```

```
## Estimate Std. Error z value
```

```
## 3|4 -418.37 67.28 -6.218
```

```
## 4|5 -415.90 67.28 -6.182
```

```
## 5|6 -412.86 67.28 -6.137
```

```
## 6|7 -410.28 67.27 -6.099
```

```
## 7|8 -408.05 67.27 -6.066
```

```
## 8|9 -404.36 67.27 -6.011
```

```
convergence(w.clm)
```

```
## nobs logLik niter max.grad cond.H logLik.Error
## 3919 -4366.67 8(0) 9.08e-10 7.6e+11 <1e-10
##
##          Estimate Std.Err Gradient Error Cor.Dec Sig.Dig
## 3|4      -4.184e+02 67.283458 6.13e-14 1.19e-09      8      11
## 4|5      -4.159e+02 67.280195 3.05e-13 1.19e-09      8      11
## 5|6      -4.129e+02 67.275900 -1.75e-11 1.19e-09      8      11
## 6|7      -4.103e+02 67.270990 1.31e-11 1.19e-09      8      11
## 7|8      -4.080e+02 67.269152 -4.41e-13 1.19e-09      8      11
## 8|9      -4.044e+02 67.270358 -1.05e-12 1.19e-09      8      11
## 'fixed acidity'      1.838e-01 0.066260 4.37e-11 -8.87e-13     11     11
## 'volatile acidity'  -5.132e+00 0.344504 1.68e-12 -9.01e-14     12     13
## 'citric acid'        1.618e-01 0.272837 1.84e-12 -2.72e-13     12     12
## 'residual sugar'     2.180e-01 0.025465 4.63e-11 -4.29e-13     12     12
## chlorides           -1.039e+00 1.528166 3.18e-13 -5.28e-12     10     11
## 'free sulfur dioxide' 1.300e-02 0.002538 1.95e-10 6.77e-15     13     12
## 'total sulfur dioxide' -7.160e-04 0.001106 9.08e-10 -5.01e-15     13     10
## density             -4.270e+02 68.174508 5.78e-12 1.20e-09      8      11
## pH                  1.642e+00 0.326757 1.71e-11 -3.96e-12     11     12
## sulphates           1.724e+00 0.291667 2.68e-12 -1.63e-12     11     12
## alcohol             4.620e-01 0.084731 4.97e-11 1.36e-12     11     11
##
## Eigen values of Hessian:
## 2.396e+07 1.660e+05 2.282e+04 1.826e+04 9.750e+02 6.943e+02 2.706e+02 1.335e+02 7.145e+01 1.818e+01
##
## Convergence message from clm:
## (3) Model is nearly unidentifiable: large eigenvalue ratio
## - Rescale variables?
## In addition: Absolute and relative convergence criteria were met
```

```
tim <- cbind(dfw.test, predict(w.clm, newdata = dfw.test, type = "class")$fit)
head(do.call("cbind", predict(w.clm, se.fit=TRUE, interval=TRUE))))
```

```
##          fit      se.fit      lwr      upr
## [1,] 0.4402663 0.01939092 0.4026815 0.4785486
## [2,] 0.2596296 0.01590626 0.2296869 0.2919962
## [3,] 0.5231015 0.01467095 0.4943021 0.5517480
## [4,] 0.5332407 0.01094017 0.5117501 0.5546087
## [5,] 0.5332407 0.01094017 0.5117501 0.5546087
## [6,] 0.5231015 0.01467095 0.4943021 0.5517480
```